

A Modern, Visual Introduction to AWS

(Beginner-Friendly Edition)

Introduction

Amazon Web Services (AWS) is a platform where you can rent computing resources—like virtual machines or storage—on demand. Traditionally, you’d buy your own physical servers and handle hardware yourself, but AWS eliminates much of that work, letting you pay only for the services you use.



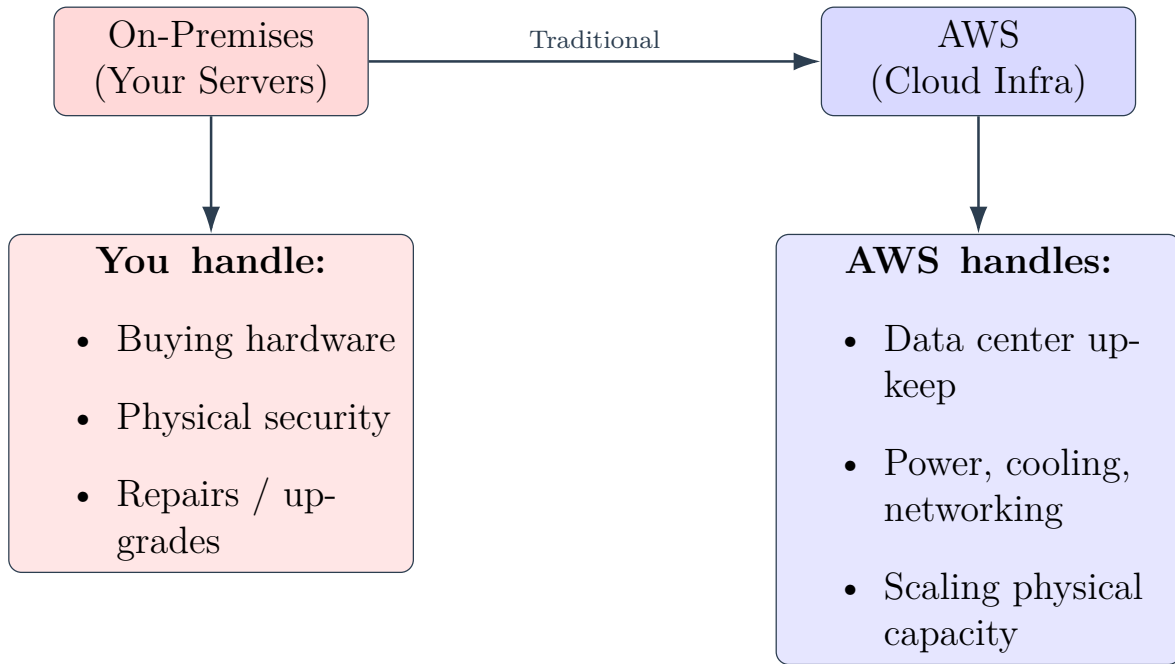
Instead of maintaining a data center, AWS handles physical security, power, cooling, and more. This gives you more time to build solutions rather than worry about hardware.

- Learn how AWS compares to traditional (on-premises) servers.
- Explore the major AWS services and how they fit together.
- See diagrams that illustrate common AWS architectures.
- Check out simple project ideas to get hands-on experience.

Let's begin!

1 Comparing On-Premises vs. AWS

1.1 High-Level Differences



On-premises means you manage everything yourself: building or renting space for servers, dealing with hardware failures, etc. AWS handles those tasks and charges you for the resources you actually use.

2 What is AWS?


2.1 Core Service Areas

AWS offers many types of services, but some core ones include:

- **Compute (EC2, Lambda)**: How you run your code or applications.
- **Storage (S3, EBS)**: Where you store files, backups, or disk volumes.
- **Databases (RDS, DynamoDB)**: Choose relational (SQL) or NoSQL data stores.
- **Networking (VPC, CloudFront)**: Control network traffic and deliver content globally.
- **Analytics (Kinesis, Athena)**: Process data streams or query large data sets.
- **Machine Learning (SageMaker)**: Tools to build and train models, then deploy them.

2.2 Regions and Availability Zones

AWS organizes its data centers worldwide into **Regions** (like US East 1, EU West 1, etc.). Each Region has multiple **Availability Zones** to improve reliability and fault tolerance.



3 How Companies Use AWS

3.1 Hosting Websites

From a simple static site stored in **S3** to a large dynamic application running on **EC2** or container services.



3.2 Data Storage and Backup

Companies keep backups, media, or data archives in **S3**, which provides high durability and easy retrieval.

3.3 Application Development

Developers build microservices using **Lambda** and **API Gateway**, then store structured data in **RDS** or **DynamoDB**.

4 Simple AWS Architectures

Below are more small diagrams that explain common patterns.

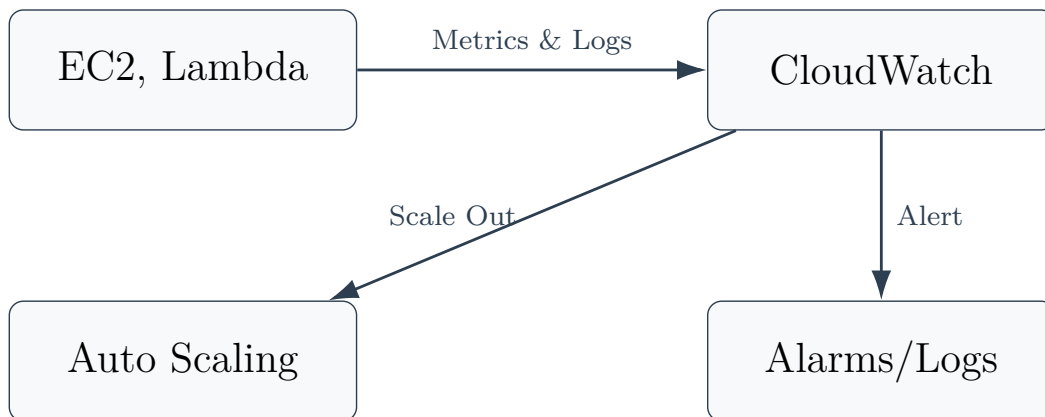
4.1 Serverless File Processor



4.2 Event-Driven Workflow



4.3 Monitoring with CloudWatch



5 Key AWS Services (Simplified)

5.1 S3 vs. EBS (Two Types of Storage)




S3 is for storing files ("objects") in buckets. **EBS** acts like a hard drive attached to an EC2 server.

5.2 EC2 vs. Lambda (Compute Options)



EC2 is a virtual machine you configure. **Lambda** runs your code in response to events without you managing servers.

5.3 Databases: RDS vs. DynamoDB



RDS supports databases like MySQL or PostgreSQL, while **DynamoDB** is a key-value NoSQL store.

5.4 CloudFront (CDN)

Provides faster delivery of static or dynamic content by using "edge locations" near users around the world.

5.5 Sample Lambda Code

```
def processImage(event, context):  
    fileKey = event["Records"][0]["s3"]["object"]["key"]  
    # Perform some image processing...  
    return {"statusCode": 200, "body": "Success"}
```

6 Hands-On Project Ideas

6.1 Beginner Projects

1. Static Website on S3/CloudFront

Host a personal or portfolio site on S3, then speed it up globally with CloudFront.

2. Simple File Processor

When you upload a file, Lambda resizes or converts it. Great for learning event-driven tasks.

6.2 Intermediate Projects

1. Full Web App (EC2, RDS, S3)

Run a web server on EC2, store user data in RDS, and keep images in S3.

2. Serverless API (API Gateway + Lambda + DynamoDB)

Everything is serverless. Lambda handles business logic, DynamoDB stores data.

6.3 Advanced Projects

1. Real-Time Data (Kinesis + Lambda + DynamoDB)

Stream logs or sensor data into Kinesis, process in Lambda, save in DynamoDB.

2. Machine Learning (SageMaker + S3 + Lambda)

Store training data in S3, train a model with SageMaker, then call it via Lambda.



7 Extra Visuals & Step-by-Step Flows

7.1 Typical User Request Flow



7.2 Simple Auto Scaling



8 A Learning Path

8.1 Step 1: Core Services

- **S3** for storing files
- **Lambda** for event-driven code
- **DynamoDB** for simple NoSQL

8.2 Step 2: Building APIs and Servers

- **API Gateway** for REST calls
- **EC2** for custom servers
- **RDS** for relational data

8.3 Step 3: Advanced Tools

- **Kinesis** for streaming data
- **SageMaker** for machine learning
- **Athena/Redshift** for analytics



Conclusion

AWS lets you run applications in the cloud instead of managing your own hardware. You can store static files in **S3**, launch code on **Lambda**, or run a full server via **EC2**. Distributing apps across Regions and Availability Zones improves reliability.

Key Takeaways:

- **Start small:** Try a static site on S3 or a simple file-processor with Lambda.
- **Add complexity:** Use RDS or DynamoDB to store data, and CloudFront to speed up content delivery.
- **Explore advanced features:** Handle real-time data (Kinesis) or train machine learning models (SageMaker).

For more details, check out the [official AWS docs](#), and use the **AWS Free Tier** to learn without large costs.