

A Simple, Visual Introduction to AWS

(Improved Edition)

January 18, 2025

Contents

Introduction	1
1 What is AWS?	2
1.1 Definition	2
1.2 Key Benefits	2
1.3 Service Categories	2
1.4 High-Level Visualization	2
2 How Companies Use AWS	3
2.1 Web Hosting	3
2.1.1 Static Hosting with S3 and CloudFront	3
2.2 Data Storage and Backup	3
2.3 Application Development	3
2.4 Machine Learning	3
2.5 Analytics	3
3 How You Can Use AWS	4
3.1 Starting Small	4
3.2 Workflow Automation Example	4
3.3 Monitoring and Logging	4
4 Key AWS Services	5
4.1 S3 (Simple Storage Service)	5
4.2 Lambda (Serverless Compute)	5
4.3 DynamoDB (NoSQL Database)	5

4.4	EC2 (Elastic Compute Cloud)	5
4.5	CloudFront (Content Delivery Network)	5
5	Suggested Projects	6
5.1	Beginner Projects	6
5.2	Intermediate Projects	6
5.3	Advanced Projects	6
6	Additional Visualizations	7
6.1	Application Architecture Example	7
6.2	Monitoring Feedback Loop	7
7	Learning Path	8
7.1	1. Start with Basics	8
7.2	2. Add More Services Gradually	8
7.3	3. Tackle Advanced Tools	8
	Conclusion	9

Introduction

This updated document aims to give newcomers an easy-to-follow, visually guided introduction to **Amazon Web Services (AWS)**. Whether you're building a small personal project or diving into large-scale deployments, AWS offers a wide range of services to help you succeed. This guide focuses on:

- **What AWS is and why it matters**
- **Common use cases in the real world**
- **Step-by-step learning projects**
- **Many diagrams to illustrate core concepts**

1 What is AWS?

1.1 Definition

Amazon Web Services (AWS) is a cloud computing platform that lets you:

- Rent computing resources on-demand
- Store, process, and analyze data
- Host and scale applications globally

Key Benefits:

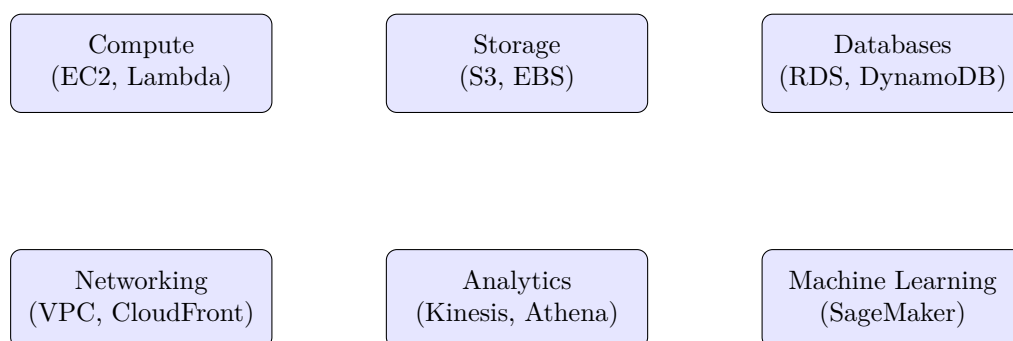
- **Pay-As-You-Go:** Only pay for the resources you actually use.
- **Scalability:** Effortlessly adjust capacity based on real-time traffic.
- **Global Infrastructure:** Deploy and deliver content across multiple regions worldwide.

1.2 Service Categories

AWS is large, but it helps to see services grouped by functionality:

- **Compute:** EC2 (virtual servers), Lambda (serverless)
- **Storage:** S3 (object storage), EBS (block storage)
- **Database:** RDS (relational), DynamoDB (NoSQL)
- **Networking:** VPC, CloudFront (CDN)
- **Analytics:** Kinesis, Athena
- **Machine Learning:** SageMaker

1.3 High-Level Visualization



2 How Companies Use AWS

2.1 Web Hosting

Many companies host their websites or web applications on AWS. Options range from:


- **Static sites:** Simple S3 + CloudFront setup
- **Dynamic sites:** EC2-based servers or container services (ECS, EKS)

2.1.1 Example: Static Hosting with S3 and CloudFront



2.2 Data Storage and Backup


Many organizations use **Amazon S3** for storing large files, media, or backups. Its highly durable (99.999999999% durability) and integrates well with other AWS services.



2.3 Application Development and Deployment

AWS supports various application architectures, from traditional server-based (EC2) to fully **serverless** (Lambda). Developers often use:

- **API Gateway** to create REST or WebSocket APIs
- **Lambda** to run code without provisioning servers
- **DynamoDB** or **RDS** as the database layer



2.4 Machine Learning

AWS offers **SageMaker** to simplify building, training, and deploying ML models. Common workflow:

- Store dataset in S3
- Train model in SageMaker
- Deploy model endpoint

2.5 Analytics

Analyze large volumes of data or real-time streams:

- **Kinesis** for real-time data ingestion
- **Athena** to query S3 data with standard SQL

3 How You Can Use AWS


3.1 Starting Small

Even beginners can build:

- **Static portfolio site:** Hosted on S3 + CloudFront
- **File processor:** Upload an image to S3, resize it using Lambda, and store the result

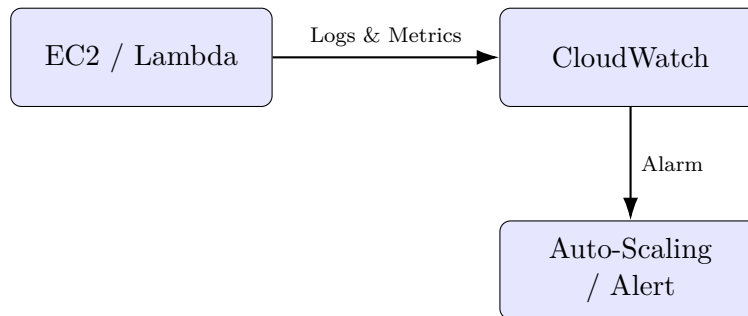
3.2 Workflow Automation

AWS excels at **event-driven** workflows:



3.3 Monitoring and Logging


CloudWatch collects logs from your AWS resources, provides metrics, and can trigger alarms or scaling actions when thresholds are met.



4 Key AWS Services (Simplified Overview)

4.1 S3 (Object Storage)

A scalable, secure, and cost-effective way to store large quantities of files.



4.2 Lambda (Serverless Compute)

Write code functions that respond to events (e.g., S3 uploads, API Gateway calls). No need to manage servers.

```
def ProcessFile(event, context):
    S3_BUCKET_NAME = "my-bucket"
    fileKey = event["Records"][0]["s3"]["object"]["key"]

    # ... do something interesting here ...

    return {
        "statusCode": 200,
        "body": "Processing complete"
    }
```

4.3 DynamoDB (NoSQL)

Key-value store optimized for performance and scalability. Excellent choice for user profiles, session data, or IoT data.

4.4 EC2 (Virtual Servers)

Allows you to rent and configure virtual machines. Ideal if you need complete control over the OS and software environment.

4.5 CloudFront (CDN)

A global content delivery network that caches and delivers content (static or dynamic) from edge locations near your users.

5 Suggested Projects for Hands-On Learning

5.1 Beginner Level

1. Static Website

- Host a personal resume or portfolio on S3
- Use CloudFront to speed up global delivery

2. File Processor with Lambda and S3

- Upload an image to S3
- Lambda automatically resizes/optimizes the image
- Outputs to another S3 bucket or the same bucket

5.2 Intermediate Level

1. Web App with EC2, RDS, and S3

- EC2: runs your backend
- RDS: relational database
- S3: static images or media

2. Serverless API (API Gateway + Lambda + DynamoDB) Create a REST or GraphQL API. Lambda handles business logic; DynamoDB stores the data.


5.3 Advanced Level

1. Real-Time Data (Kinesis + Lambda + DynamoDB)

- Kinesis streams incoming data
- Lambda processes records in near real time
- DynamoDB stores aggregated or processed data


2. Machine Learning (SageMaker + S3 + Lambda)

- S3 for training data
- SageMaker to train and host the model
- Lambda or a web app calls the deployed model endpoint

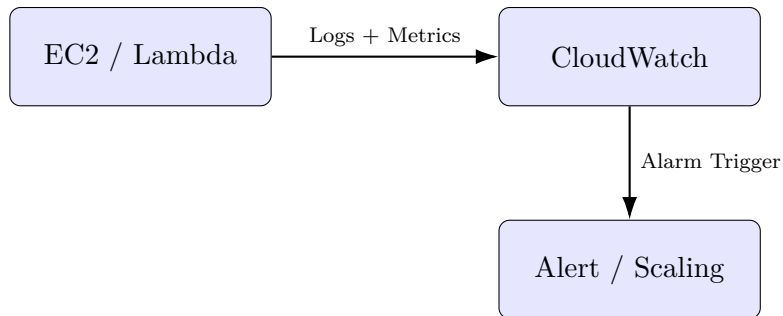


6 Additional Visualizations

6.1 Application Architecture (Serverless Example)



6.2 Monitoring and Alerting Loop



7 Learning Path

7.1 1. Foundation

Focus on:

- S3 (basic storage)
- Lambda (serverless compute)
- DynamoDB (simple NoSQL)

Build small automation or workflow projects to get comfortable with the AWS console and basic configurations.

7.2 2. Intermediate Services


Once comfortable, expand:

- **API Gateway** for building APIs
- **RDS** for relational data
- **EC2** for custom environments

7.3 3. Advanced Tools

Tackle big-data or specialized domains:

- **Kinesis** for real-time event streaming
- **SageMaker** for machine learning
- **Redshift** or **Athena** for analytics



Conclusion

We've covered the basics of **AWS** in a simple, visually guided way. By starting with fundamental services like S3, Lambda, and DynamoDB, you can quickly grasp the core principles of the AWS ecosystem. Then you can incrementally add services—EC2, RDS, API Gateway—to handle more complex or performance-intensive applications. Eventually, you can move on to real-time data processing (Kinesis), large-scale analytics (Redshift, Athena), and machine learning (SageMaker).

Throughout this journey:

- Keep experimenting with small, hands-on projects.
- Use the **AWS Free Tier** to control costs.
- Refer to official AWS Documentation for deeper detail.

Happy Building!